

Spring 2025: Neural Networks & Deep Learning – ICP -6

Assignment -6

Name: Vanitha Chintalapudi

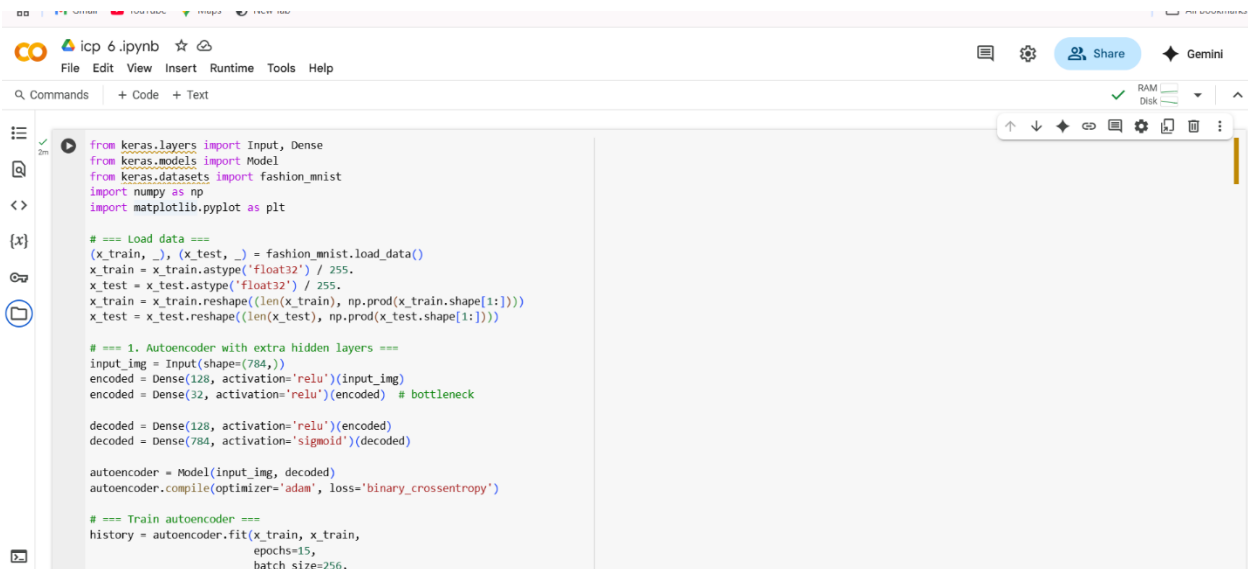
Student ID: 700756782

Github Link: <https://github.com/VanithaChintalapudi10/Neural-network-deep-learning>

Video Link:

https://drive.google.com/file/d/1OmOJQLiMaBG7ujRWbXzSZLCl6webni1/view?usp=drive_link

1. Add one more hidden layer to autoencoder
2. Do the prediction on the test data and then visualize one of the reconstructed version of that test data. Also, visualize the same test data before reconstruction using Matplotlib
3. Repeat the question 2 on the denoising autoencoder
4. plot loss and accuracy using the history object



```
from keras.layers import Input, Dense
from keras.models import Model
from keras.datasets import fashion_mnist
import numpy as np
import matplotlib.pyplot as plt

# === Load data ===
(x_train, _), (x_test, _) = fashion_mnist.load_data()
x_train = x_train.astype('float32') / 255.
x_test = x_test.astype('float32') / 255.
x_train = x_train.reshape((len(x_train), np.prod(x_train.shape[1:])))
x_test = x_test.reshape((len(x_test), np.prod(x_test.shape[1:])))

# === 1. Autoencoder with extra hidden layers ===
input_img = Input(shape=(784,))
encoded = Dense(128, activation='relu')(input_img)
encoded = Dense(32, activation='relu')(encoded) # bottleneck

decoded = Dense(128, activation='relu')(encoded)
decoded = Dense(784, activation='sigmoid')(decoded)

autoencoder = Model(input_img, decoded)
autoencoder.compile(optimizer='adam', loss='binary_crossentropy')

# === Train autoencoder ===
history = autoencoder.fit(x_train, x_train,
                        epochs=15,
                        batch_size=256,
```

colabresearch.google.com/drive/1v3-zPiDAVRPvZ3S671YJPMuiMCWN93lp#scrollTo=cv7V5EYffynk

GmailYouTubeMapsNew Tab

All Bookmarks

icp 6.ipynb

File Edit View Insert Runtime Tools Help

Commands + Code + Text

RAM Disk

2m

```
batch_size=256,
shuffle=True,
validation_data=(x_test, x_test))

# === 2. Predict & visualize ===
decoded_imgs = autoencoder.predict(x_test)

n = 5
plt.figure(figsize=(10, 4))
for i in range(n):
    # Original
    ax = plt.subplot(2, n, i + 1)
    plt.imshow(x_test[i].reshape(28, 28), cmap='gray')
    plt.title("Original")
    plt.axis('off')

    # Reconstructed
    ax = plt.subplot(2, n, i + 1 + n)
    plt.imshow(decoded_imgs[i].reshape(28, 28), cmap='gray')
    plt.title("Reconstructed")
    plt.axis('off')
plt.tight_layout()
plt.show()

# === 3. Denoising Autoencoder ===

# Add noise
noise_factor = 0.5
x_train_noisy = x_train + noise_factor * np.random.normal(loc=0.0, scale=1.0, size=x_train.shape)
x_test_noisy = x_test + noise_factor * np.random.normal(loc=0.0, scale=1.0, size=x_test.shape)
```

icp 6.ipynb

File Edit View Insert Runtime Tools Help

Commands + Code + Text

2m

```
ax = plt.subplot(3, n, i + 1)
plt.imshow(x_test_noisy[i].reshape(28, 28), cmap='gray')
plt.title("Noisy")
plt.axis('off')

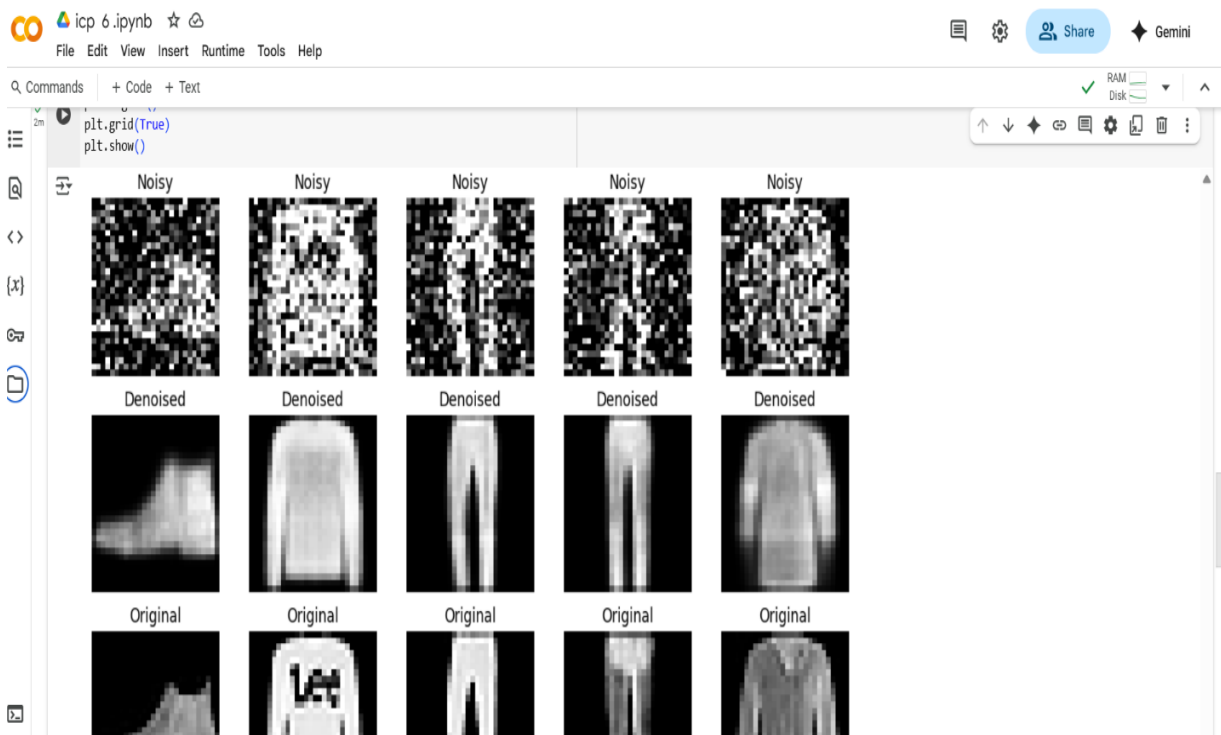
# Denoised
ax = plt.subplot(3, n, i + 1 + n)
plt.imshow(denoised_imgs[i].reshape(28, 28), cmap='gray')
plt.title("Denoised")
plt.axis('off')

# Ground truth
ax = plt.subplot(3, n, i + 1 + 2 * n)
plt.imshow(x_test[i].reshape(28, 28), cmap='gray')
plt.title("Original")
plt.axis('off')
plt.tight_layout()
plt.show()

# === 4. Plot training/validation loss ===
plt.plot(history.history['loss'], label='Train Loss - AE')
plt.plot(history.history['val_loss'], label='Val Loss - AE')
plt.plot(denoise_history.history['loss'], label='Train Loss - DAE')
plt.plot(denoise_history.history['val_loss'], label='Val Loss - DAE')
plt.title("Loss over Epochs")
plt.xlabel("Epoch")
plt.ylabel("Loss")
plt.legend()
plt.grid(True)
plt.show()
```

```
plt.grid(True)
plt.show()

Downloading data from https://storage.googleapis.com/tensorflow/tf-keras-datasets/train-labels-idx1-ubyte.gz
29515/29515 — 0s 0us/step
Downloading data from https://storage.googleapis.com/tensorflow/tf-keras-datasets/train-images-idx3-ubyte.gz
26421888/26421888 — 0s 0us/step
Downloading data from https://storage.googleapis.com/tensorflow/tf-keras-datasets/t10k-labels-idx1-ubyte.gz
5148/5148 — 0s 0us/step
Downloading data from https://storage.googleapis.com/tensorflow/tf-keras-datasets/t10k-images-idx3-ubyte.gz
4422102/4422102 — 0s 0us/step
Epoch 1/15
235/235 — 8s 25ms/step - loss: 0.4481 - val_loss: 0.3138
Epoch 2/15
235/235 — 5s 19ms/step - loss: 0.3078 - val_loss: 0.3008
Epoch 3/15
235/235 — 4s 18ms/step - loss: 0.2973 - val_loss: 0.2955
Epoch 4/15
235/235 — 3s 15ms/step - loss: 0.2920 - val_loss: 0.2918
Epoch 5/15
235/235 — 5s 14ms/step - loss: 0.2894 - val_loss: 0.2893
Epoch 6/15
235/235 — 4s 18ms/step - loss: 0.2867 - val_loss: 0.2876
Epoch 7/15
235/235 — 4s 14ms/step - loss: 0.2849 - val_loss: 0.2867
Epoch 8/15
235/235 — 5s 15ms/step - loss: 0.2840 - val_loss: 0.2858
Epoch 9/15
235/235 — 4s 19ms/step - loss: 0.2820 - val_loss: 0.2838
Epoch 10/15
235/235 — 4s 15ms/step - loss: 0.2813 - val_loss: 0.2829
Epoch 11/15
```



Q Commands

+ Code + Text

